

**Problem Statement:** Predict health insurance premiums (medical charges) using demographic, lifestyle, and medical history variables from the **Insurance\_Prediction** table (1M records).

#### Business Constraint:

- First **700k** → model training, Next **200k** → validation, Last **100k** → production/live scoring
- Same preprocessing must run during **training and prediction** and Deliverable must support **batch prediction or real-time API**

**Target:** charges (continuous → regression problem)

---

#### Available features:

- Age, Gender, BMI, Children, Smoker, Region, Occupation, Medical history, Family medical history, Exercise frequency, Coverage level

	age	gender	bmi	children	smoker	region	medical_history	family_medical_history	exercise_frequency	occupation	coverage_level	charges
0	46.0	male	21.45	5.0	yes	southeast	Diabetes	None	Never	Blue collar	Premium	20460.307669
1	25.0	female	25.38	2.0	yes	northwest	Diabetes	High blood pressure	Occasionally	White collar	Premium	20390.899218
2	38.0	male	44.88	2.0	yes	southwest	None	High blood pressure	Occasionally	Blue collar	Premium	20204.476302
3	25.0	male	19.89	0.0	no	northwest	None	Diabetes	Rarely	White collar	Standard	11789.029843
4	49.0	male	38.21	3.0	yes	northwest	Diabetes	High blood pressure	Rarely	White collar	Standard	19268.309838

---

### 3. Data Split Strategy (Business Requirement)

Dataset	Records	Purpose
Train	7,00,000	Model learning
Validation	2,00,000	Model evaluation
Production	1,00,000	Real-world prediction

```
[5]: train = df.iloc[:700000].reset_index(drop=True)
val   = df.iloc[700000:900000].reset_index(drop=True)
prod  = df.iloc[900000: ].reset_index(drop=True)

print(train.shape, val.shape, prod.shape)
```

(700000, 12) (200000, 12) (100000, 12)

---

### 4. Data Cleaning & Preprocessing

Performed inside **preprocessing.py** to guarantee consistency.

#### Missing Value Treatment

##### Numerical

- age, children → filled with median
- children rounded and cast to integer

```
[13]: #filling the value as median for age and children
for i in ['age','children']:
    median_val=train[i].median()
    train[i]=train[i].fillna(median_val)

#rounding the values in column children
train['children']=train['children'].round()

#changing the data type to astype(Int64)
train['children']=train['children'].astype('Int64')
```

## Categorical

- gender → mode
- medical\_history, family\_medical\_history → "No\_Record"
- exercise\_frequency, occupation → "Unknown"

```
[17]: # data entry issue and user skipped, Filling with mode = "male" does not change medical meaning much.  
train['gender']=train['gender'].fillna(train['gender'].mode()[0])  
  
[18]: # missing value might mean: no tests done, No diagnosis, not recorded, hence filling with No_Record  
train['medical_history']=train['medical_history'].fillna('No_Record')  
  
[19]: train['family_medical_history']=train['family_medical_history'].fillna('No_Record')  
  
[20]: # Missing value mean : Person didn't answer, Person unsure, Not recorded So missing is not equal to any existing category, hence unknown  
train['exercise_frequency']=train['exercise_frequency'].fillna('Unknown')  
  
[21]: train['occupation']=train['occupation'].fillna('Unknown')
```

---

## Encoding Strategy

Feature	Encoding Type
Gender, Smoker	Binary
exercise_frequency, coverage_level	Ordinal
medical_history, family_medical_history, Occupation	One-Hot

```
[26]: train['gender']=train['gender'].map({'male':0, 'female':1})  
[27]: train['smoker']=train['smoker'].map({'yes':1,'no':0})  
[28]: train=pd.get_dummies(train, columns=['region', 'medical_history','family_medical_history', 'occupation'], prefix=['region','med_h','fam_med_h','occ'], dropna=True)  
[29]: train['exercise_frequency']=train['exercise_frequency'].map({'never':0, 'rarely':1, 'occasionally':2, 'frequently':3})  
[30]: train['coverage_level']=train['coverage_level'].map({'basic':0, 'standard':1, 'premium':2})
```

---

## Scaling

StandardScaler applied to:

- age, bmi, children

**Purpose:** Stable coefficients, Compatible with linear and regularized models and Consistent preprocessing pipeline

---

## 5. Feature Selection Approach

### 1. Target correlation analysis

Correlation with target

```
[43]: corrX_train.join(y_train).corr()['charges'].sort_values(ascending=False)
corr

[43]: charges      1.000000
smoker       0.565662
coverage_level  0.463364
med_h_heart_disease 0.393801
fam_med_h_heart_disease 0.392399
exercise_frequency 0.164179
occ_white_collar 0.128043
bmi          0.103391
children     0.076237
age           0.060833
region_southeast 0.001091
occ_unknown    0.000546
region_northwest -0.026399
region_southwest -0.040140
occ_student    -0.062523
gender         -0.112955
occ_unemployed -0.127895
fam_med_h_high_blood_pressure -0.131265
med_h_high_blood_pressure -0.133094
med_h_no_record   -0.261450
fam_med_h_no_record -0.262272
Name: charges, dtype: float64
```

### 2. Variance threshold

```
[4]: from sklearn.feature_selection import VarianceThreshold
vt=VarianceThreshold(threshold=0.01)
vt.fit(X_train)

low_variance_features=X_train.columns[~vt.get_support()]
low_variance_features
```

```
[4]: Index([], dtype='object')
```

### 3. VIF (multicollinearity check)

LvGJ.	feature	VIF
2	bmi	10.031681
0	age	8.123823
3	children	3.079653
5	exercise_frequency	2.688370
6	coverage_level	2.409913
4	smoker	1.953795
12	med_h_no_record	1.915086
10	med_h_heart_disease	1.915067

### 4. Linear coefficients

```
coef=pd.Series(model.coef_, index=X_train.columns).sort_values(ascending=True)
coef.head()

fam_med_h_no_record      -867.128325
med_h_no_record          -865.845489
occ_unemployed            -642.432681
gender                   -488.444795
fam_med_h_high_blood_pressure -433.060098
dtype: float64
```

### 5. LASSO shrinkage

```
lcoef=pd.Series(lasso.coef_, index=X_train.columns).sort_values(ascending =True)
lcoef

fam_med_h_no_record      -866.276265
med_h_no_record          -864.985182
occ_unemployed            -641.500015
gender                   -487.945112
fam_med_h_high_blood_pressure -432.170124
med_h_high_blood_pressure -431.738917
occ_student                -426.509287
region_southwest           -343.666968
region_northwest            -300.538433
region_southeast             -213.653191
occ_unknown                 -103.626456
occ_white_collar            215.646838
age                        271.113425
children                  336.090186
bmi                      461.519182
```

### 6. Random Forest feature importance

```
[57]: importance=pd.Series(model.feature_importances_, index=X_train.columns).sort_values(ascending=True)
importance
```

occ_unknown	0.000614
med_h_high blood pressure	0.006365
fam_med_h_high blood pressure	0.006438
children	0.006661
age	0.007866
occ_student	0.008810
occ_white collar	0.009351
gender	0.012150
occ_unemployed	0.012419
bmi	0.016950
med_h_no_record	0.019050
fam_med_h_no_record	0.019466
exercise_frequency	0.028074
fam_med_h_heart_disease	0.153693
med_h_heart_disease	0.155016

## Consensus Finding

Strong predictors:

- Smoker, Coverage level, Medical history, Family medical history, Exercise frequency

Weak but retained:

- Gender, Children, Occupation

**Reason:** Some models still used them and no multicollinearity risk

**Region dropped** since consistently weak

---

## Models Selection

Model	R <sup>2</sup>	Training Time	Inference	Model Size
Linear Regression	~0.985	<b>0.31 sec</b>	<b>0.01 sec</b>	<b>~0.8 KB</b>
Lasso	~0.985	0.37 sec	0.006 sec	0.75 KB
Random Forest	~0.986	98 sec	8.39 sec	<b>~GBs</b>

## Hence, Final Model Decision: Lasso (Production-friendly)

- Similar accuracy to Random Forest, 10,000× smaller, Real-time API feasible, Interpretable coefficients, Low infra cost
- 

## Final Pipeline architecture

### 1. Train.py

- Load Regression.db
- Business Split (700k train / 200k val)
- Calls preprocess on both train and validation data
- Train Final Model -> lasso
- Evaluate (R2, RMSE)
- Save **model.pkl + scaler.pkl + feature\_columns.json**

```
(base) C:\Users\shambhavic\OneDrive - Temenos\Data Science and Analytics\jupyteruploads\Capstone Project\Regression\Insurance premium prediction>python train.py
Data loaded: (1000000, 12)
(700000, 12) (200000, 12) (100000, 12)
Validation R2: 0.9878
Validation RMSE: 562.77
Model and scaler saved.
```

## 2. Preprocessing.py

- Copy dataframe
- Handle missing values
- Text standardization (strip + lower)
- Encoding: Binary, Ordinal, One-hot
- Drop region
- Separate target if present
- Scale numeric columns
- Return X, y, scaler

## 3. Predict.py

- Load model.pkl, scaler.pkl and feature\_columns.json
- Load production data (last 100k)
- preprocess(prod\_df, scaler, training=False)
- Align columns with training schema
- Predict premiums
- Save predicted\_premiums.csv

```
model and scaler saved.
(base) C:\Users\shambhavic\OneDrive - Temenos\Data Science and Analytics\jupyteruploads\Capstone Project\Regression\Insurance premium prediction>python predict.py
Data loaded: (1000000, 12)
Production data shape: (100000, 12)
   predicted_charges
0          17619.217455
1          9935.562123
2          14560.510963
3          16563.611803
4          16359.428511
Predictions saved to predicted_premiums.csv
```

A	B	C	D	E	F	G	H	I	J	K	L	M	N	O
1	age	gender	bmi	children	smoker	region	medical_history	family_medical_history	exercise_frequency	occupation	coverage_level	charges	predicted_charges	
2	64	male	34.67	2	no	northwest	Heart disease	Heart disease	Never	Student	Basic	17685.37836	17619.21746	
3	18	male	40.19	5	no	southwest	High blood pressure		Rarely	Blue collar	Basic	9547.613014	9935.562123	
4	27	female	46.74	4	no	northeast	High blood pressure	Diabetes	Never	Student	Premium	14692.00132	14560.51096	

## 4. App.py

- User JSON inputConvert to DataFrame
- preprocess(input\_df, scaler, training=False)
- Align columns
- model.predict()
- Return premium JSON

```
Predictions saved to predicted_premiums.csv
(base) C:\Users\shambhavic\OneDrive - Temenos\Data Science and Analytics\jupyteruploads\Capstone Project\Regression\Insurance premium prediction>python app.py
 * Serving Flask app 'app'
 * Debug mode: on
WARNING: This is a development server. Do not use it in a production deployment. Use a production WSGI server instead.
 * Running on http://127.0.0.1:5000
Press CTRL+C to quit
 * Restarting with watchdog (windowsapi)
 * Debugger is active!
 * Debugger PIN: 357-612-689
127.0.0.1 - - [08/Feb/2026 18:39:46] "GET / HTTP/1.1" 200 -
127.0.0.1 - - [08/Feb/2026 18:40:30] "POST /predict HTTP/1.1" 200 -
127.0.0.1 - - [08/Feb/2026 18:43:01] "POST /predict HTTP/1.1" 200 -
 * Detected change in 'C:\Users\shambhavic\AppData\Local\anaconda3\Scripts\conda-script.py', reloading
 * Restarting with watchdog (windowsapi)
 * Debugger is active!
 * Debugger PIN: 357-612-689
```



POST http://127.0.0.1:5000/

http://127.0.0.1:5000/predict

POST http://127.0.0.1:5000/predict

Params Authorization Headers (8) Body  Pre-request Script Tests Settings

None  form-data  x-www-form-urlencoded  raw  binary  JSON

1 {  
2 "age": 45,  
3 "gender": "male",  
4 "bmi": 28.5,  
5 "children": 2,  
6 "smoker": "no",  
7 "region": "northeast",  
8 "medical\_history": "diabetes",  
9 "family\_medical\_history": "no\_record",  
10 "exercise\_frequency": "occasionally",  
11 "occupation": "student",  
12 "coverage\_level": "basic"  
13 }  
14  
15

Body Cookies Headers (5) Test Results

Pretty Raw Preview Visualize JSON

Status: 200 OK Time: 142 ms Size: 210 B Save Response

1 "predicted\_premium": 12488.39242837531

POST http://127.0.0.1:5000/

http://127.0.0.1:5000/predict

POST http://127.0.0.1:5000/predict

Params Authorization Headers (8) Body  Pre-request Script Tests Settings

None  form-data  x-www-form-urlencoded  raw  binary  JSON

1 {  
2 "age": 45,  
3 "gender": "male",  
4 "bmi": 28.5,  
5 "children": 2,  
6 "smoker": "no",  
7 "region": "northwest",  
8 "medical\_history": "diabetes",  
9 "family\_medical\_history": "no\_record",  
10 "exercise\_frequency": "never",  
11 "occupation": "student",  
12 "coverage\_level": "basic"  
13 }  
14  
15

Body Cookies Headers (5) Test Results

Pretty Raw Preview Visualize JSON

Status: 200 OK Time: 112 ms Size: 211 B Save Response

1 "predicted\_premium": 11769.894499545268